

Inf2C Computer Systems

Coursework 3

MIPS Cache Simulator

Deadline: Friday, 11 December, 16:00

Instructor: Aaron Smith

TA: David Schall

In this assignment you will extend the simulator from Coursework 2 for a 5-stage multi-cycle MIPS processor with three different cache models. You will be provided with working skeleton code for a MIPS processor simulator to extend with your cache implementations. You are strongly advised to read up on caches in the lecture notes and course textbook and to commence work as soon as possible.

This is the third and last assignment for the Inf2C-CS course and is worth 20% of the overall course mark. Please bear in mind that the guidelines on academic misconduct from the Undergraduate year 2 student handbook are available at the following link <http://web.inf.ed.ac.uk/infweb/student-services/ito/students/year2>.

1 Overview

In this assignment you will extend the provided code to support three different cache models between the processor and the memory. By reading from input files, the skeleton feeds the memory and register state file with instructions and data, and simulates the standard 5-stage instruction execution cycle. The output functions are already implemented and called in the skeleton code.

Skeleton Organization: The skeleton is broken down into four source code files. Each file accomplishes particular tasks. You are allowed to modify only certain files. The functionality and permission to modify for each file are described in Table 1.

mipssim.c: This file describes the multi-cycle MIPS processor as studied in class. The processor consists of the following core components: PC, Pipeline registers (IR, A,

Filename	Functionality	Modifiable
<code>mipssim.c</code>	Multi-cycle MIPS processor (datapath + control)	No
<code>mipssim.h</code>	Data structure definitions for datapath	No
<code>memory_hierarchy.c</code>	Memory hierarchy implementation	Yes
<code>parser.h</code>	Reading and parsing input files	No

Table 1: Source Files

B, MDR, and ALUOut), Programmer-visible registers (in the register file), ALU, ALU Control, and Control.

The processor’s functionality can be broken down into the following logical stages: *instruction_fetch*, *decode_and_read_RF*, *execute*, *memory_access* and *write_back*. Note that each stage updates some architectural or microarchitectural state. For instance, the *instruction_fetch* stage updates the IR. The *write_back* stage updates the RF.

Furthermore, this file handles the Control component by implementing a finite state machine. The state machine can be found in a function named *FSM*.

mipssim.h: This file defines the following required data structures for the MIPS processor:

- *ctrl_signals*, which control the datapath and are updated on a cycle-by-cycle basis by the control FSM
- *instr_meta*, which stores information about the instruction currently stored in the IR
- *memory_stats_t*, which consists of memory stats for loads, stores and instruction fetches
- *pipe_regs*, which includes the PC and microarchitectural registers of the processor (IR, A, B, ALUOut and MDR). For convenience, we refer to these registers as *pipeline registers* to indicate that they are spread out over the datapath (unlike the programmer-visible registers, which are all located inside the register file).

On any given cycle, the complete state of the processor is stored in a structure called *architectural_state*. This structure includes the current clock cycle since the start of execution (*clock_cycle*), state of the current instruction’s execution (e.g., INSTR.FETCH or DECODE), current values of control signals (*control*) and memory stats (*mem_stats*). This structure also includes an array of *registers*, which models a register file, as well as the *memory*. Finally, *architectural_state* also includes the pipeline registers, which are maintained in two *pipe_regs* structs: *curr_pipe_regs* and *next_pipe_regs*. The former (*curr_pipe_regs*) is used within a cycle to *read* the value of a given register. Meanwhile, a value that needs to be written into a pipeline register at the end of the cycle should be stored in *next_pipe_regs*. At the end of each clock cycle, *curr_pipe_regs* is updated with values from *next_pipe_regs*; this functionality is already provided for you.

memory_hierarchy.c: This file provides the memory interface via two functions: *memory_read*, which is used for reading from memory and *memory_write*, which is used

2.2 Task 2: Fully Associative Cache

In Task 2 you will implement a fully associative cache, which means that data can be stored in any cache block. Your implementation should use a least recently used (LRU) cache replacement policy. Use the lowest index in case LRU determines more than one set.

2.3 Task 3: 2-way Set Associative Cache

In Task 3 you will implement a 2-way set associative cache. For this cache the index is used to find the set and the tag is used to find the cache block within the set. Your implementation should use a least recently used (LRU) cache replacement policy. Use the lowest index in case LRU determines more than one set.

2.4 Cache Implementation

You must implement the caches in the `memory_hierarchy.c` file. In this file, there are two memory functions already defined: `memory_read` and `memory_write`. The `memory_read` function is used to fetch instructions and load data from memory. The `memory_write` function writes data to the memory. You must extend these functions to access the cache. You must also update the relevant cache hits statistic, as well as the content of the cache, as necessary. The statistics are maintained in a global structure `arch_state.mem_stats`, which is updated on every access to the memory hierarchy.

The caches will use the following parameters:

Fixed parameters:

1. Addresses are 32-bits
2. Memory is byte addressable
3. Cache block size is 16 bytes
4. Cache holds both instructions and data (this is called a *unified* cache)
5. Cache uses a write-through policy
6. Cache uses a write-no-allocate policy; i.e., if a store does not find the target block in the cache, it does not allocate it.
7. LRU replacement policy (when applicable)

Variable parameters:

1. The size of the cache (in bytes) is passed as the first command-line argument to the program and stored for you in a global variable named `cache_size`. Note that this parameter specifies the size of the *data* portion of the cache. Any other information that the cache needs to store (e.g., tags) are additional to this.

2. The type of the cache is passed as the second command-line argument to the program and stored for you in another global variable named *cache_type*. The type is given as integer value and defined as 1: direct mapped, 2: fully associative, 3: 2-way set associative.

Notes:

- Instruction fetch uses the *memory_read* function. Hence, each instruction fetch updates the cache hit statistics just like loads.
- If the cache size is set to **0**, the cache is disabled and the processor directly accesses main memory.
- The cache will not be larger than 16KB.
- You **must** use dynamic memory allocation (*malloc*) for your cache data structure.

Based on the cache size, you need to calculate the number of bits required for the tag field, which will be stored in *arch_state.bits_for_cache_tag*. You must set this variable by providing a correct value in the call to the *memory_stats_init* function in *memory_hierarchy.c*. The function itself is already defined for you.

Summary: In this assignment, you are required to implement three different cache models in a MIPS simulator. You will need to create and configure the appropriate cache data structure(s) based on parameters that are passed as command-line arguments. You will need to set the *arch_state.bits_for_cache_tag* variable to the number of bits for the tag via a call to *memory_stats_init* and update the cache hit statistics (in *arch_state.mem_stats*) as necessary for both instruction fetches and data accesses.

3 Notes on the Implementation

1. You are expected to dynamically allocate memory using **malloc** for the cache data structure. Submissions that use statically allocated memory (e.g., statically declared arrays of fixed size) will be penalized. Variable-length arrays that can be sized based on a runtime parameter in C99 are **NOT** allowed.
2. There are many ways to implement the cache structure. The details of the implementation are entirely up to you – you are the designer! The only thing that matters is correctness at the functional level. Just remember, you are coding in a high level language – think data structures, not bits and gates.
3. You can use the C library functions available from the header files that are already included in *mipssim.c* and *mipssim.h*.
4. You will not be marked on how fast your simulator runs or how much memory it uses. The only criteria for evaluating your implementation of the caches are correct.

4 Input Files

The skeleton will read two files:

1. *The memory state file*: a text file, which is a mix of instructions and data represented by a sequence of 0 and 1 characters, one word (32 bits) per line. Note that the *first* non-comment line of the memory file is always an instruction, which starts at address 0x0. Subsequent words are placed in consecutive memory locations. For this assignment, a special instruction with opcode *111111* (in binary) is considered as the End-Of-Program (EOP) instruction, which terminates the program. In the provided skeleton, there are two memory state file examples: `memfile-simple.txt` and `memfile-complex.txt`.

2. *The register state file*: a text file, which contains the initial state of programmer-visible registers. This file has up to 31 uncommented lines for registers (i.e., all of the programmer-visible registers except \$0) starting with register \$1. Each line specifies a decimal value to which the corresponding register will be initialized. If fewer than 31 values are specified, the remaining registers will be initialized with zeros. In the provided skeleton, there is one register state file example: `regfile.txt`.

5 Output Format

You do not need to generate any output for marking purposes. Instead, you must ensure that in each cycle all relevant variables/structs of the *arch_state* struct are properly updated. The automated marking will check the *arch_state* struct in each cycle, checking the correctness of the control signals (i.e., the *arch_state.control* field), and datapath state (i.e., fields: *arch_state.curr_pipe_regs* and *arch_state.registers*). The automated marking will use the functions *marking_after_clock_cycle()* and *marking_at_the_end()*. Make sure that you **do not use or modify** these functions!

6 Debugging

1. To verify correctness of your implementation, you should write your own memory and register state files. For verifying the cache, you need to come up with test cases to predictably generate certain behaviors (hits and misses). For instance, think of a trace with four accesses that uses four different addresses and has a 50% cache hit rate in a direct-mapped cache.
2. Your simulator will be tested with memory and register state files different from the provided ones.
3. You can use *printf* function for debugging your code. It will not affect your marking.

7 Compiling and Running the Simulator

You **must** compile the simulator on the DICE machines with the following command:

```
gcc -o mipssim mipssim.c memory_hierarchy.c -std=gnu99 -lm
```

Note that this is the exact command we will use for compiling your code for marking purposes. Compiling the source files creates an executable `mipssim`. Make sure that your simulator both compiles with the exact command and runs on a **DICE** machine without errors and warnings. Otherwise, you will receive a **0** mark.

The following are examples of invoking the simulator with valid command-line parameters.

```
./mipssim 1024 1 memfile-simple.txt regfile.txt
```

Where `mipssim` is the name of the executable file. `1024` indicates cache is enabled and set to a size of 1024 bytes. `1` configures the *cache type* as direct mapped (2 is fully associative and 3 is 2-way set associative). Additionally, `memfile-simple.txt` is the name of the memory file and `regfile.txt` is the name of the register state file. Both memory and register state files are located in the same directory as `mipssim`.

8 Submission

The skeleton code and example inputs are contained in the course work Git repository on GitLab: <https://git.ecdf.ed.ac.uk/asmith47/inf2c-cs-20/cw3>

You will modify the skeleton files provided for the tasks and push your changes to your fork of the coursework Git repo on GitLab. You can push changes to GitLab any time up until the submission deadline. We will not mark any changes pushed after the submission deadline has passed.

The submission deadline is 4pm on December 11, 2020.

For additional information about late penalties and extension requests, see the School web page below. Do **not** email any course staff directly about extension requests – we are unable to grant them. Instead, follow the instructions on the web page: <http://web.inf.ed.ac.uk/infweb/student-services/ito/admin/coursework-projects/late-coursework-extension-requests>

9 Assessment

The assignment will be auto-marked. Your solutions will be evaluated with a number of test inputs. For each task, your mark will be proportional to the pass rate of the tests.

10 Similarity Checking and Academic Misconduct

You must submit your own work. Any code that is not written by you must be clearly identified and explained through comments at the top of your files. Failure to do so is plagiarism. Detailed guidelines on what constitutes plagiarism can be found at: <http://web.inf.ed.ac.uk/infweb/admin/policies/guidelines-plagiarism>. All submitted code is checked for similarity with other submissions using the MOSS ¹ system. MOSS has been effective in the past at finding similarities. It is not fooled by name changes and reordering of code blocks.

11 Questions

If you have any questions about the assignment, please start by checking existing discussions on Piazza – chances are, others have already encountered (and, possibly, solved) the same problem. If you can't find the answer to your question, start a new discussion. You should also take advantage of the drop-in labs and the lab demonstrators who are there to answer your questions.

November 24, 2020

¹<http://theory.stanford.edu/~aiken/moss/>